

Comprehensive Assessment of Proposed Doctoral Research in Frugal AI and Small Language Models

The landscape of artificial intelligence is currently undergoing a massive structural transformation. For the past several years, the prevailing philosophy in the field was that larger models inherently yielded better performance. **This led to an arms race resulting in massive Large Language Models (LLMs) containing hundreds of billions, or even trillions, of parameters. While these models possess remarkable capabilities, their sheer scale has created unsustainable economic, environmental, and infrastructural bottlenecks. Training and deploying these massive systems require exorbitant computational resources, massive server farms, and intense energy and water consumption, fundamentally limiting their accessibility to only the largest technology conglomerates and excluding billions from the digital economy.**

In response to this severe resource strain, the academic and industrial focus has rapidly pivoted toward "Frugal AI" and Small Language Models (SLMs). **Frugal AI advocates for using the minimum level of complexity required to effectively solve a problem.** It promotes solutions that are sustainable, cost-effective, and capable of operating in resource-constrained environments, such as edge devices, mobile phones, and on-premises hardware without continuous cloud connectivity. **By utilizing advanced techniques such as model compression, quantization, pruning, and targeted fine-tuning, SLMs can capture the core reasoning capabilities of larger models while requiring a fraction of the computational power.**

The six research papers proposed for this doctoral thesis form a highly cohesive, forward-looking, and socially relevant thematic arc. They traverse the highly technical mechanics of model compression, the behavioral economics of managerial adoption, the sociology of corporate strategy, and the democratization of digital resources for smaller economic players. The foundational premise of the research is incredibly strong.

However, because Frugal AI is currently one of the most aggressively researched domains globally, several of the foundational questions posed in this initial research outline have already been addressed by the scientific community, particularly in studies published between 2024 and early 2026. If the research simply asks questions that have already been answered, the resulting thesis will resemble a literature review rather than a novel, defensible doctoral contribution.

To ensure the resulting doctoral thesis offers a groundbreaking contribution to the field, this report exhaustively evaluates each proposed paper against the current state-of-the-art literature. The language used in this report is intentionally accessible, breaking down complex technical jargon into clear, understandable concepts. For each topic, the analysis identifies

what has already been established by existing research, uncovers the underlying mechanisms driving these phenomena, and provides highly specific recommendations for pivoting the research track to guarantee academic novelty and rigorous, original contributions.

Paper 1: Impact of Model Quantization on Algorithmic Bias

The initial premise of this paper asks a crucial question: [do algorithms become biased if we quantize them or reduce their size?](#) To understand this, [one must first understand quantization.](#) In simple terms, [quantization is the process of reducing the numerical precision of a model's underlying weights.](#) Think of it like compressing a high-resolution photograph to save storage space; the overall image remains recognizable, but the fine, granular details become blurred or pixelated. [In AI, weights are often reduced from highly precise 16-bit floating-point numbers to much smaller 8-bit or 4-bit integers.](#)

Current State of the Research

The short answer to the foundational question is yes; quantization significantly and unpredictably alters the social and demographic bias of a model. [However, this question has been comprehensively answered by recent literature, meaning an exploratory paper will lack novelty.](#)

[Extensive studies conducted between 2024 and 2026 have meticulously mapped the exact effects of model compression on social biases.](#)¹ Recent large-scale evaluations, including studies analyzing over five million responses across numerous bias benchmarks, have established that the effects of quantization extend far beyond mere performance degradation.⁴ [Quantized models exhibit unpredictable shifts in social bias that can completely undo the safety alignments and behavioral guardrails installed by developers during the training process.](#)

[Specifically, research demonstrates that aggressive compression—such as moving to 4-bit weights—tends to slightly increase stereotypes and unfairness in generative tasks.](#)³ For instance, when models are asked to complete sentences or make decisions involving different demographic groups, quantized models are much more likely to revert to historical stereotypes.⁵ Interestingly, the impact of quantization is not universally negative. [While stereotyping increases, quantization has also been shown to inadvertently reduce the model's tendency to generate highly toxic content and neutralize its overall sentiment, creating a highly complex and contradictory ethical landscape.](#)³

A critical and fascinating discovery in recent literature is a phenomenon termed "quantization-induced masked bias flipping".⁴ Studies indicate that up to 21% of a model's responses can flip between biased and unbiased states after the model is compressed.⁴ Because some responses flip from being biased to unbiased, while others flip from unbiased to

biased, the *overall average* bias score of the model might look exactly the same as the original, uncompressed model. This average score is incredibly misleading because it masks severe, underlying volatility and harm directed at specific demographic groups.⁴ Furthermore, the type of acceleration strategy used matters immensely. Weight quantization (methods like AWQ or GPTQ) has been shown to cause drastic, unpredictable impacts on bias, whereas compressing the model's memory cache (KV-cache quantization) remains relatively robust and preserves the base model's ethical alignment.²

Second and Third-Order Insights

To produce a doctoral-level thesis, one must understand the hidden mechanisms driving these changes. The core issue is the loss of representational granularity. **During the later stages of training (known as alignment or fine-tuning), models learn highly specific, delicate mathematical weights to suppress biased outputs and adhere to safety guidelines.¹ When these numbers are truncated and rounded off during quantization, the delicate mathematical boundaries that separate a neutral, safe response from a biased, harmful one are eroded.** Because the compressed model loses the mathematical precision required to navigate complex social nuances, it defaults to the broader, often biased statistical patterns hidden deep within its original internet training data.¹

The broader implication here is highly concerning for the AI industry: current safety alignment techniques are largely superficial. **If simply shrinking a file size reintroduces harmful behaviors that were supposedly "aligned" away, it indicates that current safety training does not fundamentally fix the model's worldview. It merely places a fragile, high-precision mathematical filter over a fundamentally biased base model.**

Recommended Research Pivot

Since the core question—does quantization cause bias?—has been definitively answered, continuing on the original track will result in redundant research. The track must shift from simply *identifying* the problem to *engineering a mitigation strategy*.

Original Proposal Focus	Current Literature Status	Recommended PhD Pivot Strategy
Do algorithms become biased if we quantize them?	Resolved. Yes, quantization causes unpredictable bias shifts, increases stereotyping, and induces masked bias flipping. ³	Bias-Aware Quantization: Develop a novel quantization algorithm that mathematically penalizes precision loss in the specific neural pathways responsible for ethical

		alignment.
--	--	------------

Instead of measuring the bias, the research should propose a technical framework for "Ethical Quantization." As a PhD researcher, the goal should be to identify the specific parts of the neural network (e.g., specific attention heads or parameter clusters) that actively govern safety and demographic neutrality. Once these areas are mapped, the thesis could propose a "mixed-precision" quantization strategy. In this approach, the "safety-critical" parts of the brain are kept at a high resolution (8-bit or 16-bit), while the non-critical parts of the brain are aggressively compressed to 2-bit or 4-bit. This would maintain the speed and cost efficiency of Frugal AI while mathematically guaranteeing the preservation of the model's ethical alignment.

Paper 2: Managerial Decision-Making in the Selection of AI Models

The proposed research seeks to understand the trade-offs managers make between efficiency and ethics or safety when selecting AI models. The proposed methodology is highly practical: utilizing a simulated dashboard presenting different AI models with varying costs, processing speeds, and safety ratings to see if managers are willing to sacrifice safety to save money.

Current State of the Research

The integration of AI into managerial decision-making is a rapidly expanding field of study, particularly concerning how AI influences the allocation of resources, strategic planning, and the creation of sustainable business models.⁸ Existing literature heavily discusses the theoretical tension between AI efficiency and ethical constraints.¹⁰ For instance, studies have explored the "alignment problem" in business processes, highlighting how the drive for operational efficiency and lower costs can inadvertently lead to discriminatory hiring practices, unfair resource allocation, or skewed financial analysis if ethical frameworks are ignored.¹¹

Recent empirical studies have begun to examine the psychology of how humans allocate trust and decision-making power to AI agents. Utilizing frameworks like the Socio-Cognitive Model of Trust, researchers have found that managers generally do not want to hand over total control. In managerial tasks, individuals tend to assign approximately 25% to 30% of the decision-making weight to AI agents, retaining the vast majority of the authority for themselves.¹⁴ Furthermore, studies show that when managers are confronted with AI systems that appear to possess high autonomy or "free will," their willingness to collaborate actually decreases. This is likely due to an "uncanny valley effect," where managers feel their human agency and authority are threatened by the machine.¹⁴

While the theoretical trade-off between Environmental, Social, and Governance (ESG) criteria and AI efficiency is well-documented⁸, behavioral studies that calculate the exact monetary value managers assign to AI safety remain remarkably scarce. Researchers note that improper

management of AI tools, especially complex dashboards, can lead to cognitive overload and "alert fatigue," ultimately reducing decision efficiency.¹⁵

Second and Third-Order Insights

The proposed dashboard experiment is exceptional because it touches on a profound psychological and economic mechanism: bounded rationality in the face of algorithmic opacity. When managers are presented with a dashboard, they are forced to make a discrete choice. In traditional business procurement, speed and cost are tangible, easily quantifiable metrics with direct, immediate impacts on the quarterly bottom line. Conversely, "safety" or "ethical alignment" is an abstract, probabilistic, and long-term risk metric.¹⁰

The underlying insight is that unless AI safety is tied to strict regulatory compliance fines or quantifiable liability costs, rational economic actors (managers) will systematically underprice safety. The phenomenon of "ethical drift" occurs when organizations gradually and imperceptibly alter their moral reasoning to favor the high-efficiency, low-cost outputs of AI, slowly eroding human oversight to maximize savings.¹⁷

A further implication is the concept of "responsibility laundering." Managers might consciously select highly efficient, lower-safety models, operating under the assumption that if an algorithmic failure occurs (such as a racially biased hiring decision or an unsafe medical recommendation), the blame can be deferred to the "black box" technology rather than personal managerial negligence.¹⁶ The dashboard experiment perfectly captures the moment this ethical compromise occurs.

Recommended Research Pivot

This proposed track is exceptionally strong and highly novel. While the theoretical background is established in the literature, the proposed empirical methodology—using a simulated dashboard to force procurement decisions—is a brilliant application of behavioral economics to AI governance.

Original Proposal Focus	Current Literature Status	Recommended PhD Pivot Strategy
Do managers sacrifice safety to save money when selecting AI models via a dashboard?	Partially Addressed. Theoretical trade-offs are widely known ¹¹ , but empirical quantification of managerial behavior is severely lacking.	Retain and Expand. Execute the dashboard experiment as a formal discrete choice experiment to calculate the exact financial elasticity of AI

		ethics.
--	--	---------

To elevate this paper to a rigorous doctoral level, the experiment should employ conjoint analysis—a statistical technique used in market research to determine how people value different attributes that make up an individual product or service. This will allow the researcher to calculate the exact "willingness to pay" (WTP) for AI safety.

By manipulating the dashboard variables across different simulated industries (e.g., comparing a high-stakes environment like healthcare resource allocation ¹¹ against a low-stakes environment like retail marketing), the research can pinpoint the exact dollar threshold at which a manager decides the cost savings of a frugal AI model outweigh the ethical risks. Furthermore, the experiment could introduce regulatory variables into the dashboard. By simulating the presence of severe regulatory fines (like those proposed by the EU AI Act), the research could definitively prove how government policy influences the managerial pricing of algorithmic ethics, providing highly valuable data for both business leaders and policymakers.

Paper 3: Governing the Edge - A Framework for Deploying Frugal AI Responsibly

This paper aims to develop a risk framework dictating when compressed, frugal models are acceptable for use (for example, drafting internal emails) and when they are strictly forbidden (for example, screening resumes for hiring or making medical diagnoses).

Current State of the Research

The governance of AI systems based on risk categorization is not a new concept; in fact, it is the foundational premise of contemporary global AI policy. The most prominent and widely discussed example is the European Union's AI Act, which enforces a stratified, systemic risk framework that explicitly categorizes AI use cases into unacceptable, high, limited, and minimal risk tiers.¹⁸ Similarly, in specific sectors like healthcare, frameworks such as FAIR-AI and guidelines from the International Medical Device Regulators Forum (IMDRF) provide detailed structures for categorizing AI solutions into low, moderate, or high-risk domains, which then determines the level of human oversight required.²⁰ These frameworks universally agree that high-stakes environments require maximum human validation, while low-stakes tasks can be automated.²²

Simultaneously, the deployment of "Frugal AI" at the edge is accelerating rapidly. Driven by the urgent need to reduce the massive environmental footprint of large data centers ²³ and to enable offline, real-time capabilities on mobile phones and internet-of-things (IoT) devices, Frugal AI relies heavily on deploying compressed models directly onto consumer hardware.²⁵ However, running AI on edge devices introduces unique vulnerabilities, such as compromised security protocols, distributed denial-of-service risks, and the threat of physical tampering with

edge hardware.²⁷

Second and Third-Order Insights

The intersection of Frugal AI and risk governance reveals a massive, unaddressed gap in existing policy. Current regulatory frameworks, like the EU AI Act, assess risk based almost entirely on the *application domain*—meaning they look at *how* the AI is being used (e.g., medical diagnostics versus administrative sorting).²⁸ However, these policies completely fail to account for the *architectural degradation* inherent in Frugal AI.

When a large language model is heavily compressed to fit onto a smartphone or an edge device, it suffers from "catastrophic forgetting" of edge-case knowledge, a reduction in logical reasoning capacity, and, as established in the analysis of Paper 1, an unpredictable degradation of its safety guardrails.¹ Therefore, a theoretically "low-risk" application—such as summarizing a simple customer service transcript—can spontaneously transform into a high-risk liability if the highly compressed edge model suddenly hallucinates toxic content, leaks private data, or generates discriminatory text because its safety weights were pruned away to save space.

Furthermore, Frugal AI fundamentally decouples the model from centralized oversight. Large frontier models accessed via the cloud can be monitored, audited, and patched centrally by the developer. An SLM deployed locally on millions of decentralized edge devices cannot be easily updated or recalled. If an edge model begins exhibiting bias, the organizational liability is highly distributed and incredibly difficult to contain.³¹

Recommended Research Pivot

The idea of creating a risk framework is conceptually sound, but proposing a framework based solely on high-stakes versus low-stakes environments will overlap heavily with existing regulatory documents like the EU AI Act. To make this a unique PhD contribution, the framework must specifically address the unique technical deficiencies of *compressed* models.

Original Proposal Focus	Current Literature Status	Recommended PhD Pivot Strategy
A Risk Framework for deploying Frugal AI based on high vs. low stakes use cases.	Heavily Researched. High vs. low-stakes risk frameworks are the standard in global policy (e.g., EU AI Act, IMDRF). ¹⁹	Compression-Specific Governance: Create a dynamic framework that dictates deployment based on the mathematical degradation index of the compressed model.

The research should evolve from a generic risk policy into a "Frugal AI Degradation Matrix." This framework would argue that a model's suitability for a task should be governed by mapping the environmental stakes directly against the model's physical compression ratio. For example, the thesis could propose specific mathematical thresholds: if a model undergoes more than a 50% sparsity pruning, or is quantized down to 2-bit precision, it is automatically and mathematically disqualified from any task involving autonomous decisions about human beings, regardless of the industry. This moves the paper away from a subjective policy discussion into an objective, sociotechnical governance model that bridges computer science and corporate compliance.

Paper 4: Algorithm Isomorphism and the Homogenization of Competitive Strategy

This paper asks a deeply sociological business question: whether the widespread adoption of LLMs for decision-making, marketing, and strategy will homogenize corporate behavior, effectively eroding the unique differences that allow companies to compete.

Current State of the Research

The concept of AI driving organizational homogenization is not only highly relevant but has already been formally recognized and termed "Algorithmic Isomorphism" in recent academic literature.³³ Drawing heavily on Neo-Institutional Theory—specifically the foundational work by DiMaggio and Powell on how organizations copy each other to gain legitimacy—researchers have posited that as organizations increasingly rely on deep learning algorithms for decision-making, they will inevitably converge on identical operational and strategic behaviors.³⁶

Recent scholarship in 2024 and 2025 has explicitly examined this phenomenon in the context of Large Language Models. Papers with titles like "ChatGPT is incredible (at being average)" analyze the frequently overlooked issue of output homogenization, describing how generative AI models structure outputs in a highly recognizable, standardized manner.³⁴ Because these models are designed to optimize for the statistical mean of their training data, they inherently produce "average" strategies.³³ It is actively recognized in management literature that while using AI improves the baseline efficiency and creativity of an individual struggling firm, it severely reduces the collective creative diversity of the broader market, trapping companies in what researchers are calling a "Great AI Convergence".³⁹

Second and Third-Order Insights

The mechanism driving Algorithmic Isomorphism is deeply rooted in the fundamental architecture of LLMs. Large language models are, at their core, advanced probability engines designed to predict the most likely next sequence of words based on massive historical datasets. By definition, the "most likely" sequence is the most average, conventional, and

historically established sequence.

When multiple competing firms use the exact same foundation models (e.g., OpenAI's GPT-4, Anthropic's Claude) to generate their marketing copy, optimize their supply chains, or forecast market trends, they are all drinking from the exact same statistical well.³⁹ According to classic business strategy, true competitive advantage relies on distinctiveness—performing different activities from rivals or performing similar activities in vastly different ways.³⁹ Algorithmic Isomorphism actively destroys this distinctiveness by ironing out the contrarian beliefs, unique corporate cultures, and human eccentricities that make a brand unique.

A critical third-order consequence of this homogenization is the heightened vulnerability of global markets to systemic shocks. In biological ecosystems, monocultures (growing only one type of crop) are highly susceptible to disease because a single pathogen can eradicate the entire uniform population. Similarly, if an entire sector's strategic decision-making is homogenized by the same algorithmic logic, an external economic shock or an embedded algorithmic hallucination could trigger simultaneous, synchronized failures across multiple competing firms. This means a localized software error could escalate into a systemic market crisis.

Recommended Research Pivot

Because the theoretical concept of "Algorithmic Isomorphism" regarding LLMs has already been coined and actively published over the last two years³⁵, a purely theoretical paper arguing that this *might* happen will lack academic novelty. The PhD thesis must pivot from theoretical philosophy to empirical measurement.

Original Proposal Focus	Current Literature Status	Recommended PhD Pivot Strategy
Do LLMs homogenize competitive strategy (Algorithm Isomorphism)?	Theorized and Published. The term is actively used, and the theory linking LLMs to strategic convergence is established. ³⁵	Empirical Quantification: Conduct a longitudinal, data-driven analysis to mathematically prove and measure the rate of strategic convergence across an industry.

To make a groundbreaking contribution, the research should employ natural language processing (NLP) to analyze real-world corporate data. The candidate could scrape publicly available documents—such as 10-K financial filings, investor call transcripts, and external marketing materials—of Fortune 500 companies in a specific sector over a five-year period encompassing the AI boom (e.g., 2021 to 2026). By using text embeddings and mathematical

cosine similarity metrics, the paper could quantitatively demonstrate whether the linguistic and strategic distance between competing firms has mathematically narrowed since the mass adoption of generative AI. Proving the theory with hard data would be a massive contribution to management science.

Paper 5: Model Compression and the Erasure of Cultural Nuance

This proposed paper questions whether the act of model compression disproportionately erases cultural nuances, prunes minority languages, and ignores local norms in favor of dominant linguistic structures.

Current State of the Research

The disparity in LLM performance across different languages is a well-documented and highly discussed challenge. Standard LLMs are overwhelmingly trained on English-centric, Western data, leading to inherent struggles with low-resource languages and non-Western cultural contexts.⁴²

Crucially, the impact of model compression (such as pruning, quantization, and knowledge distillation) on these already fragile multilingual capabilities is currently a major focal point of 2025 and 2026 NLP research.⁴⁵ Recent comprehensive evaluations of quantized models across dozens of languages reveal that compression techniques disproportionately affect "long-tail" linguistic features. In statistics, the "long tail" refers to data points that occur infrequently; in language modeling, these long-tail features inherently correspond to underrepresented languages, regional dialects, and unique cultural nuances.⁴⁵

For instance, studies evaluating model pruning in multilingual contexts show that compressing models leads to a significant degradation in capturing socio-cultural contexts embedded in indigenous or low-resource languages.⁴⁶ Furthermore, research has found that safety filters and alignment techniques applied to compressed models often exhibit a strong "Western bias," effectively erasing minority perspectives in contexts such as humor, political dialogue, or cultural norms, thereby enforcing a homogenized, hegemonic viewpoint.⁴⁹

Second and Third-Order Insights

The erasure of cultural nuance during compression is not an accidental bug; it is the fundamental mathematical objective of the compression process. Techniques like magnitude-based pruning identify and eliminate the neural weights that contribute the least to the overall predictive accuracy of the model.⁵¹ Because minority languages, regional slang, and specific local customs appear so infrequently in the global training corpus, the mathematical parameters encoding these concepts have very small magnitudes. Consequently, when an algorithm is instructed to prune a model to achieve 50% smaller file size, it identifies these vital

cultural nuances as "redundant" or "unimportant" and systematically deletes them to preserve the dominant (usually English/Western) linguistic structures.⁵⁰

The third-order implication of this erasure is a phenomenon akin to "digital cultural colonialism." As Frugal AI and SLMs become the standard for deployment in developing economies and the Global South (due to local hardware, power, and internet constraints), the very models deployed to these populations will be the ones that have been mathematically stripped of local cultural contexts. Instead of democratizing AI, aggressive and blind model compression threatens to impose a highly standardized, culturally sterile version of intelligence on global populations, accelerating digital language death and alienating local users.⁴⁹

Recommended Research Pivot

The premise that compression erases minority languages and cultural nuances is already scientifically validated and documented.⁴⁵ Therefore, the research must advance beyond proving that erasure occurs and focus entirely on innovative preservation strategies.

Original Proposal Focus	Current Literature Status	Recommended PhD Pivot Strategy
Does model compression erase cultural nuance and minority languages?	Validated. Research confirms long-tail tokens and minority norms are disproportionately targeted by pruning. ⁴⁵	Culturally-Aware Pruning: Develop and test a novel compression architecture that explicitly shields localized parameter clusters from deletion.

The PhD research should focus on engineering "Culturally-Preservative Compression." The candidate could propose a modified structured pruning methodology. By mapping the specific neurons or attention heads that activate when processing minority languages or local cultural norms, the research could introduce a "cultural importance weighting" into the loss function during the pruning phase. The objective would be to prove that a Small Language Model can be successfully compressed by 50% while artificially protecting these long-tail parameters, resulting in a frugal model that retains its cultural dexterity and respect for local norms.

Paper 6: Mitigating SME Resource Disadvantages with SLMs

The final paper investigates whether the adoption of Small Language Models (SLMs) and open-weight architectures, combined with advanced prompting techniques, can successfully mitigate the massive resource disadvantages faced by Small and Medium Enterprises (SMEs) in

the digital economy.

Current State of the Research

The narrative that SMEs face significant, often insurmountable hurdles in AI adoption due to infrastructural, financial, and skill-based resource constraints is universally accepted in current business literature.⁵³ Concurrently, the emergence of highly capable SLMs—such as Microsoft's Phi family, Meta's Llama 1B/3B, and Mistral's 3B/8B models—is universally heralded by both industry analysts and academia as the definitive solution to this exact problem.⁵⁶

These models, typically possessing between 1 billion to 14 billion parameters, provide high efficiency and specialized task capabilities at an absolute fraction of the computational cost of frontier LLMs.⁵⁸ Research confirms that SLMs enable SMEs to run models locally (edge computing), which entirely mitigates their concerns regarding sensitive data privacy, expensive vendor lock-in, and unpredictable cloud API expenses.⁵⁷ Furthermore, empirical studies published in 2025 have already demonstrated that AI adoption—specifically tailored, smaller AI solutions—significantly enhances both the sustainable and economic performance of SMEs.⁶¹

Second and Third-Order Insights

The excitement surrounding SLMs for SMEs is highly justified, but this excitement obscures significant, secondary challenges that businesses face on the ground. While SLMs lower the barrier to entry regarding inference costs and data privacy, they do not inherently solve the organizational "knowledge gap." SLMs trade generalized, broad reasoning for highly specialized efficiency.⁶⁰ Out of the box, a small model is prone to higher hallucination rates in broad domains and struggles with complex, multi-step logic compared to massive frontier LLMs.⁶⁰

To bridge this capability gap and make the SLM actually useful, SMEs must rely on techniques like Retrieval-Augmented Generation (RAG), Parameter-Efficient Fine-Tuning (PEFT), and sophisticated prompt engineering. However, building a robust local RAG pipeline or executing fine-tuning requires highly specialized data engineering talent—a resource that SMEs fundamentally lack.²⁶ Thus, the resource disadvantage is not truly mitigated; it is merely shifted from hardware and cloud computing costs to specialized human capital costs.

A third-order trend emerging in the industry is the development of "Agentic AI" ecosystems powered entirely by SLMs. Rather than relying on one massive, expensive LLM to do everything, the future points toward orchestrating multiple, specialized SLMs working in tandem (multi-agent systems) to complete complex business workflows.⁶² If SMEs can leverage open-source agentic frameworks, they could potentially replicate enterprise-level operational capacity without requiring enterprise-level budgets or massive server infrastructure.

Recommended Research Pivot

Because the foundational idea—"SLMs help SMEs save money and resources"—is already an

established industry consensus, treating it as an exploratory hypothesis will result in an unoriginal thesis. The research must shift from asking *if* it works to conducting an empirical evaluation of *how well* specific implementation architectures perform.

Original Proposal Focus	Current Literature Status	Recommended PhD Pivot Strategy
Can SLMs mitigate the resource disadvantage of SMEs?	Industry Consensus. It is widely accepted and documented that SLMs lower barriers for SMEs. ⁵⁷	Empirical Capability Parity Analysis: A quantitative study comparing the real-world performance of an optimized SLM+RAG pipeline against a frontier LLM API.

The research should evolve into a highly rigorous, technical benchmarking study. The candidate could set up a specific, common SME use case (for example, legal document review, inventory forecasting, or automated customer support) and build two distinct technological pipelines to solve it.

The first pipeline would be a "heavily resourced" approach using an expensive, massive cloud API (e.g., GPT-4o or Claude 3.5 Sonnet).

The second pipeline would be a "frugal" approach using an open-weight SLM (e.g., Llama 3 3B) enhanced with advanced prompt engineering and a localized RAG architecture.

By measuring the exact differential in output accuracy, system latency, monthly infrastructure cost, and the human setup time required, the thesis would provide a definitive, empirical blueprint proving exactly *how much* of the resource disadvantage can be mitigated, and pinpointing the specific operational thresholds where the SLM approach fails.

Synthesized Conclusion for the Doctoral Candidate

The proposed research trajectory into Small Language Models and Frugal AI is highly relevant, forward-looking, and addresses some of the most critical socio-technical bottlenecks in the current artificial intelligence landscape. The candidate has successfully identified the most pressing tension points in modern AI deployment: the battle between operational efficiency versus ethics, the friction between globalized standardization versus cultural localization, and the divide between corporate dominance versus SME democratization.

However, because this specific domain is evolving at a breakneck pace, many of the

exploratory questions posed in the initial proposal have already been answered by literature published between 2024 and early 2026. Asking "does quantization cause bias?" or "do SLMs help SMEs?" will yield literature reviews, not novel PhD contributions.

To ensure this thesis yields groundbreaking, publishable research that advances the frontier of the field, the candidate must execute a strategic pivot across all six papers: **transition away from identifying phenomena, and move toward engineering mitigations and executing empirical quantifications.**

1. Move from simply measuring bias to actively building **bias-resistant quantization algorithms.**
2. Advance from theorizing ethical trade-offs to calculating the **economic elasticity of AI safety** through interactive dashboard simulations and conjoint analysis.
3. Shift from broad, generic risk policies to formulating a strict **mathematical degradation framework** for edge deployment.
4. Graduate from theorizing algorithmic isomorphism to **mining real-world corporate data** to mathematically prove strategic convergence.
5. Pivot from observing cultural erasure to engineering **culturally-preservative pruning techniques** that protect long-tail data.
6. Move beyond validating SLM cost-efficiency to **empirically benchmarking SLM+RAG architectures** against frontier models for specific SME workflows.

By adopting these specific structural pivots, the resulting doctoral thesis will evolve from a passive commentary on the current state of Frugal AI into a foundational, highly technical body of work that actively shapes how efficient, ethical, and culturally aligned AI systems are engineered, governed, and deployed in the future.

Works cited

1. Uncertainty Drives Social Bias Changes in Quantized Large Language Models - arXiv.org, accessed on March 2, 2026, <https://arxiv.org/html/2602.06181v1>
2. The Impact of Inference Acceleration on Bias of LLMs - ACL Anthology, accessed on March 2, 2026, <https://aclanthology.org/2025.naacl-long.91v1.pdf>
3. [2508.18088] How Quantization Shapes Bias in Large Language Models - arXiv.org, accessed on March 2, 2026, <https://arxiv.org/abs/2508.18088>
4. (PDF) Uncertainty Drives Social Bias Changes in Quantized Large Language Models, accessed on March 2, 2026, https://www.researchgate.net/publication/400583950_Uncertainty_Drives_Social_Bias_Changes_in_Quantized_Large_Language_Models
5. How Quantization Shapes Bias in Large Language Models - arXiv, accessed on March 2, 2026, <https://arxiv.org/pdf/2508.18088?>
6. How Quantization Shapes Bias in Large Language Models - GitHub Pages, accessed on March 2, 2026, <https://insait-institute.github.io/quantization-affects-social-bias/>
7. The Impact of Inference Acceleration on Bias of LLMs - arXiv, accessed on March

- 2, 2026, <https://arxiv.org/html/2410.22118v3>
8. Artificial Intelligence in Managerial Decision-Making for Sustainable Business Models: A Systematic Literature Review - MDPI, accessed on March 2, 2026, <https://www.mdpi.com/2079-8954/14/3/245>
 9. Assessing the Role of Generative AI and Small and Large Language Models in Enhancing Decision-Making and Sustainability in Supply Chain Management - Diva-Portal.org, accessed on March 2, 2026, <https://www.diva-portal.org/smash/get/diva2:1984938/FULLTEXT01.pdf>
 10. Ethics of Automated Decision-Making, accessed on March 2, 2026, <https://thedecisionlab.com/reference-guide/philosophy/ethics-of-automated-decision-making>
 11. Ethical implications of AI-driven clinical decision support systems on healthcare resource allocation - PMC, accessed on March 2, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11662436/>
 12. AI and management: navigating the alignment problem for ethical and effective decision-making - International Association for Computer Information Systems, accessed on March 2, 2026, https://www.iacis.org/iis/2024/2_iis_2024_488-498.pdf
 13. Bias in Decision-Making for AI's Ethical Dilemmas: A Comparative Study of ChatGPT and Claude - arXiv, accessed on March 2, 2026, <https://arxiv.org/html/2501.10484v1>
 14. Trust and AI weight: human-AI collaboration in organizational management decision-making, accessed on March 2, 2026, <https://www.frontiersin.org/journals/organizational-psychology/articles/10.3389/fo rgp.2025.1419403/full>
 15. (PDF) HOW INTERACTIVE DASHBOARDS IMPROVE MANAGERIAL DECISION-MAKING IN OPERATIONS MANAGEMENT - ResearchGate, accessed on March 2, 2026, https://www.researchgate.net/publication/389560981_HOW_INTERACTIVE_DASHBOARDS_IMPROVE_MANAGERIAL_DECISION-MAKING_IN_OPERATIONS_MANAGEMENT
 16. Ethical Risk Factors and Mechanisms in Artificial Intelligence Decision Making - PMC - NIH, accessed on March 2, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9495402/>
 17. (PDF) From Decision Support to Decision Influence: Understanding the Ethical Drift of AI in Management Practices - ResearchGate, accessed on March 2, 2026, https://www.researchgate.net/publication/398889250_From_Decision_Support_to_Decision_Influence_Understanding_the_Ethical_Drift_of_AI_in_Management_Practices
 18. Public Consultation on National Implementation of EU Harmonised Rules on Artificial Intelligence (AI Act): Submissions received, accessed on March 2, 2026, <https://enterprise.gov.ie/en/consultations/consultations-files/public-consultation-on-ai-act-submissions-received.pdf>
 19. Frontiers in Artificial Intelligence and Applications - Trepo, accessed on March 2, 2026,

- https://trepo.tuni.fi/bitstream/handle/10024/231932/Information_Modelling_and_Knowledge_Bases_XXXV.pdf?sequence=-1&isAllowed=y
20. A practical framework for appropriate implementation and review of artificial intelligence (FAIR-AI) in healthcare - ResearchGate, accessed on March 2, 2026, https://www.researchgate.net/publication/394449779_A_practical_framework_for_appropriate_implementation_and_review_of_artificial_intelligence_FAIR-AI_in_healthcare
 21. Towards Risk-Free Trustworthy Artificial Intelligence: Significance and Requirements - OPUS at UTS, accessed on March 2, 2026, <https://opus.lib.uts.edu.au/bitstream/10453/179813/2/International%20Journal%20of%20Intelligent%20Systems%20-%202023%20-%20Alzubaidi%20-%20Towards%20Risk%E2%80%90Free%20Trustworthy%20Artificial%20Intelligence%20.pdf>
 22. AI Standards for Global Impact, accessed on March 2, 2026, <https://s41721.pcdn.co/wp-content/uploads/2025/09/ai-standards-for-global-impact.pdf>
 23. Regenerate - 2050NOW La Maison, accessed on March 2, 2026, <https://2050nowlamaison.com/wp-content/uploads/2025/07/Regenerate-the-future-by-2050NOW-La-Maison.pdf>
 24. Scaling India's Data Centre Ecosystem | Systemiq, accessed on March 2, 2026, <https://www.systemiq.earth/wp-content/uploads/2026/02/Scaling-Indias-Data-Centre-Ecosystem.pdf>
 25. magazine - INSIDE Industry Association, accessed on March 2, 2026, <https://inside-association.eu/wp-content/uploads/2025/11/INSIDE-Magazine-11.pdf>
 26. Frugal AI - Iterate.ai, accessed on March 2, 2026, <https://iterate.ai/ai-glossary/frugal-ai>
 27. 2025 Accelerating a Frugal AI Ecosystem WhitePaper, accessed on March 2, 2026, <https://frugalai.org/wp-content/uploads/2025/11/2025-Nov-Accelerating-a-Frugal-AI-Ecosystem-WhitePaper.pdf>
 28. A Pragmatic Framework for Federated Learning Risk and Governance in Academic Medical Centers - JMIR AI, accessed on March 2, 2026, <https://ai.jmir.org/2026/1/e80022>
 29. TRUSTWORTHINESS FOR AI IN DEFENCE, accessed on March 2, 2026, <https://eda.europa.eu/docs/default-source/brochures/taid-white-paper-final-09052025.pdf>
 30. The rise of Large Language Models: from fundamentals to application - Management Solutions, accessed on March 2, 2026, <https://www.managementsolutions.com/sites/default/files/minisite/static/72b0015f-39c9-4a52-ba63-872c115bfbd0/llm/pdf/rise-of-llm.pdf>
 31. How to Implement an AI Governance Framework Using Safe, Ethical and Reliable AI Guardrails - Ivanti, accessed on March 2, 2026, <https://www.ivanti.com/blog/ai-governance-framework-responsible-ai-guardrails>
 32. (PDF) Inverted Automation: Generative AI, Dexterous Robotics, and the Future of Work in a Manual-Labor Society How “desk work” becomes cheap, “hands work” becomes scarce, and the social contract is forced to adapt - ResearchGate,

- accessed on March 2, 2026,
https://www.researchgate.net/publication/398934242_Inverted_Automation_Generative_AI_Dexterous_Robotics_and_the_Future_of_Work_in_a_Manual-Labor_Society_How_desk_work_becomes_cheap_hands_work_becomes_scarce_and_the_social_contract_is_forced_to_adapt
33. (PDF) ChatGPT is incredible (at being average) - ResearchGate, accessed on March 2, 2026,
https://www.researchgate.net/publication/393796279_ChatGPT_is_incredible_at_being_average
 34. The Even Sheen of AI: Kitsch, LLMs, and Homogeneity - arXiv, accessed on March 2, 2026, <https://arxiv.org/pdf/2509.16794>
 35. Algorithmic Isomorphism: the Implications of Artificial Intelligence for, accessed on March 2, 2026,
<https://merit.url.edu/ca/publications/algorithmic-isomorphism-the-implications-of-artificial-intelligen/>
 36. Institutionalization, Coercive Isomorphism, and the Homogeneity of Strategy, accessed on March 2, 2026,
<https://journals.sfu.ca/abr/index.php/abr/article/view/49>
 37. New Institutional Theory and AI: Algorithms as a Driving Force of Organizational Isomorphism - ResearchGate, accessed on March 2, 2026,
https://www.researchgate.net/publication/372854773_New_Institutional_Theory_and_AI_Algorithms_as_a_Driving_Force_of_Organizational_Isomorphism
 38. (PDF) A Review and a Proposal About Socio-economic Impacts of Artificial Intelligence, accessed on March 2, 2026,
https://www.researchgate.net/publication/345958963_A_Review_and_a_Proposal_About_Socio-economic_Impacts_of_Artificial_Intelligence
 39. How AI is creating a crisis of business sameness - RALI, accessed on March 2, 2026,
<https://ralionline.com/newsinsights/how-ai-is-creating-a-crisis-of-business-sameness/>
 40. Artificial Intelligence and Strategic Decision-Making: Evidence from Entrepreneurs and Investors | Strategy Science - PubsOnLine - INFORMS.org, accessed on March 2, 2026,
<https://pubsonline.informs.org/doi/10.1287/stsc.2024.0190>
 41. Making Choices Rational: The Elective Affinity of Artificial Intelligence and Organizational Decision-Making - Weizenbaum Library, accessed on March 2, 2026,
<https://www.weizenbaum-library.de/bitstreams/a868a0e9-6e32-4c4c-955e-a7a9dee772c3/download>
 42. CCD-Bench: Probing Cultural Conflict in Large Language Model Decision-Making, accessed on March 2, 2026,
https://www.researchgate.net/publication/396249806_CCD-Bench_Probing_Cultural_Conflict_in_Large_Language_Model_Decision-Making
 43. Seminar Report | PDF | Artificial Intelligence | Intelligence (AI) & Semantics - Scribd, accessed on March 2, 2026,

- <https://www.scribd.com/document/953234931/Seminar-Report>
44. Equitable Evaluation via Elicitation This work was supported in part by the Hire Aspirations Institute, grant G-2017-9890 from the Alfred P. Sloan Foundation, and the Simons Collaboration on Theory of Algorithmic Fairness. Authors are listed alphabetically. - arXiv.org, accessed on March 2, 2026, <https://arxiv.org/html/2602.21327v1>
 45. Quantization and Machine Translation: When Do LLMs Forget Languages? - arXiv, accessed on March 2, 2026, <https://arxiv.org/html/2508.20893v1>
 46. The 14th International Joint Conference on Natural Language Processing & The 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics - ACL Anthology, accessed on March 2, 2026, <https://aclanthology.org/events/ijcnlp-2025/>
 47. Targeted injection attack toward the semantic layer of large language models - Frontiers, accessed on March 2, 2026, <https://www.frontiersin.org/journals/computer-science/articles/10.3389/fcomp.2025.1683495/full>
 48. Asian Chapter of the Association for Computational Linguistics (2025) - ACL Anthology, accessed on March 2, 2026, <https://aclanthology.org/events/aac1-2025/>
 49. Position Paper: Losing our Tail – Again: On (Un)Natural Selection And Multilingual Large Language Models - arXiv, accessed on March 2, 2026, <https://arxiv.org/html/2507.03933v1>
 50. Large Language Models for Subjective Language Understanding: A Survey - arXiv, accessed on March 2, 2026, <https://arxiv.org/html/2508.07959v1>
 51. A Comprehensive Review: Model Compression for Large Language Models (LLMs) | by Doil Kim | Medium, accessed on March 2, 2026, <https://medium.com/@kimdoil1211/a-comprehensive-review-model-compression-for-large-language-models-llms-2c0c106d5dbe>
 52. Efficient Multi-Lingual LLM Deployment for Low-Resource Languages - ResearchGate, accessed on March 2, 2026, https://www.researchgate.net/publication/394431935_Efficient_Multi-Lingual_LLM_Deployment_for_Low-Resource_Languages
 53. The new normal: The status quo of AI adoption in SMEs - Taylor & Francis, accessed on March 2, 2026, <https://www.tandfonline.com/doi/full/10.1080/00472778.2024.2379999>
 54. Artificial Intelligence Adoption in SMEs: Survey Based on TOE–DOI Framework, Primary Methodology and Challenges - MDPI, accessed on March 2, 2026, <https://www.mdpi.com/2076-3417/15/12/6465>
 55. (PDF) Digital Transformation Strategies and Challenges in Small and Medium Enterprises (SMEs): A Systematic Review and Future Directions - ResearchGate, accessed on March 2, 2026, https://www.researchgate.net/publication/392512122_Digital_Transformation_Strategies_and_Challenges_in_Small_and_Medium_Enterprises_SMEs_A_Systematic_Review_and_Future_Directions
 56. Small Language Model: A Practical Guide for SMEs and Enterprises, accessed on March 2, 2026, <https://www.gdsonline.tech/small-language-model/>

57. What is a small language model and how can businesses leverage this AI tool? - The World Economic Forum, accessed on March 2, 2026, <https://www.weforum.org/stories/2025/01/ai-small-language-models/>
58. SCBX AI Outlook 2025: Beaconing the Future of Artificial Intelligence, accessed on March 2, 2026, https://www.scbx.com/wp-content/uploads/2025/05/SCBX-AI-Outlook-2025_EN_G.pdf
59. On-Device LLMs for SMEs: Challenges and Opportunities - arXiv.org, accessed on March 2, 2026, <https://arxiv.org/html/2410.16070v2>
60. Small Language Models and the Role They'll Play in 2025 - Software Mind, accessed on March 2, 2026, <https://softwaremind.com/blog/small-language-models-and-the-role-theyll-play-in-2025/>
61. A SEM-ANN analysis to examine impact of artificial intelligence technologies on sustainable performance of SMEs - PMC, accessed on March 2, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11825937/>
62. DIGITAL ECONOMY TRENDS 2025, accessed on March 2, 2026, <https://dco.org/wp-content/uploads/2024/12/Digital-Economy-Trends-2025.pdf>
63. AI Privacy Risks & Mitigations – Large Language Models (LLMs) - European Data Protection Board, accessed on March 2, 2026, <https://www.edpb.europa.eu/system/files/2025-04/ai-privacy-risks-and-mitigations-in-llms.pdf>